

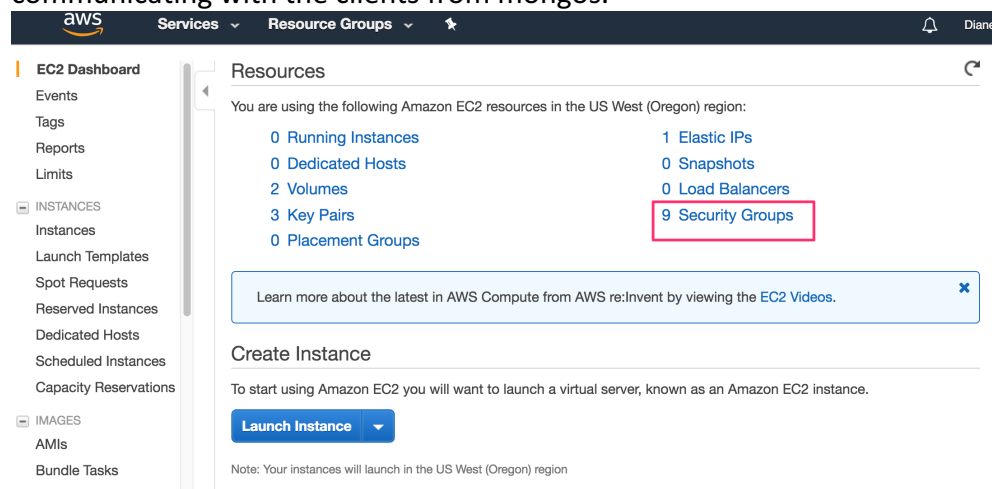
MSDS 697 Creating replicated sharded MongoDBs on AWS

Objective

Create ten servers for replication and sharding - one server for the mongos and three servers each for the first replica set, the second replica set, and the config server replica set on AWS.

Steps

Step 1. Create two security groups on EC2 → Security Groups – one for communicating between shards, replicated nodes, configuration server and mongos and the other for communicating with the clients from mongos.



1.1) Create internal access security group

port 27017 - mongos

port 27018 - shards

port 27019 - config server

Once the security group is created, change the source as its security group ID.

<input checked="" type="checkbox"/>	mongo_internal_access	sg-0d9835b0cd734277a	mongo_internal_access	vpc-c0c30ba7	mongo_internal_access
<input type="checkbox"/>		sg-69a50211	default	vpc-3e1da259	default VPC security group
<input type="checkbox"/>		sg-74a4950d	default	vpc-c0c30ba7	default VPC security group
<input type="checkbox"/>		sg-7842be0b	launch-wizard-22	vpc-c0c30ba7	launch-wizard-22 created 2018-08-10.

Description	Inbound	Outbound	Tags
-------------	---------	----------	------

Edit

Type ⓘ	Protocol ⓘ	Port Range ⓘ	Source ⓘ	Description ⓘ
Custom TCP Rule	TCP	27017 - 27019	sg-0d9835b0cd734277a	mongo_internal_acc...

1.2) Create external access security group

Create Security Group

Security group name ⓘ

Description ⓘ

VPC ⓘ

Security group rules:

Inbound Outbound

Type ⓘ	Protocol ⓘ	Port Range ⓘ	Source ⓘ	Description ⓘ
Custom TCP Rule	TCP	27017	Anywhere	0.0.0.0, ::/0

Add Rule

Cancel Create

27107 – Clients to mongos.

Step 2. Launch 10 instances (1 for mongos, 3 for configuration, 3 for shard1, and 3 for shard2)

- 2.1) Choose an Amazon Linux (Storage optimized would be a better choice but \$\$).
- a. Make sure to choose default and mongo_internal_access. (For mongos, also choose mongo_external_access.)

- 2.2) ssh to each instance and install MongoDB.

- a. Follow steps mentioned in <https://docs.mongodb.com/manual/tutorial/install-mongodb-on-amazon/>

i. Create a /etc/yum.repos.d/mongodb-org-4.0.repo and edit.

ii. `sudo yum install -y mongodb-org`

iii. start mongod (**Except for the mongos server**)

`sudo service mongod start`

Run `sudo yum update` to apply all updates.

`[[ec2-user@ip-172-31-31-46 ~]$ sudo service mongod start`

Starting mongod (via systemctl):

[OK]

- 2.3) Create an AMI to create multiple instances with the same configuration and repeat.

- 2.4) For each instance (**Except for mongos**)

- a. Create /data/db :

`sudo mkdir -p /data/db`

`[[ec2-user@ip-172-31-31-46 ~]$ sudo mkdir -p /data/db`

- b. Change the owner of /data/db to mongod (**Except for mongos**) :

`sudo chown -R mongod:mongod /data/db`

`[[ec2-user@ip-172-31-31-46 ~]$ sudo chown -R mongod:mongod /data/db`

Step 3. Create 3-node replica sets for 2 shards.

3.1) On /etc/mongod.conf of each node, update the configuration.

- Change
storage:
dbPath: /data/db
- Add port: 27018
- Add bindIp: 0.0.0.0
- Add
replication:
replSetName: rs1 (or rs2)
- Add
sharding:
clusterRole: shardsvr

where to write logging data.

```
systemLog:
  destination: file
  logAppend: true
  path: /var/log/mongodb/mongod.log
```

Where and how to store data.

```
storage:
  dbPath: /data/db
  journal:
    enabled: true
```

engine:

mmapv1:

wiredTiger:

how the process runs

```
processManagement:
  fork: true # fork and run in background
  pidFilePath: /var/run/mongodb/mongod.pid # location of pidfile
  timeZoneInfo: /usr/share/zoneinfo
```

network interfaces

```
net:
  port: 27018
  bindIp: 0.0.0.0 # Enter 0.0.0.0,:: to bind to all IPv4 and IPv6 addr
```

#security:

#operationProfiling:

```
replication:
  replSetName: rs2
```

```
sharding:
  clusterRole: shardsvr
```

3.2) Start mongod on each node using

```
sudo mongod --config /etc/mongod.conf
```

```
[ec2-user@ip-172-31-31-46 ~]$ sudo mongod --config /etc/mongod.conf
```

3.3) Connect a mongo shell to one of the mongod instances and run rs.initiate() per shard.

```
sudo mongo --port 27018
```

```
config = {_id : "rs1 (or rs2)",
```

```
members:[{_id:0,host:"PRIVATE_IP_ADDRESS:27018"},
```

```
{_id:1,host:"PRIVATE_IP_ADDRESS: 27018"},
```

```
{_id:2,host:"PRIVATE_IP_ADDRESS: 27018"}]}}
```

```
rs.initiate(config)
```

```
> config = {_id : "rs2", members:[{_id:0,host:"172.31.33.99:27018"},{_id:1,host:"172.31.45.208:27018"},{_id:2,host:"172.31.27.49:27018"}]}
```

```
{
  "_id" : "rs2",
  "members" : [
    {
      "_id" : 0,
      "host" : "172.31.33.99:27018"
    },
    {
      "_id" : 1,
      "host" : "172.31.45.208:27018"
    },
    {
      "_id" : 2,
      "host" : "172.31.27.49:27018"
    }
  ]
}
```

```
> rs.initiate(config)
```

```
{
  "ok" : 1,
  "operationTime" : Timestamp(1546827166, 1),
  "$clusterTime" : {
    "clusterTime" : Timestamp(1546827166, 1),
    "signature" : {
      "hash" : BinData(0,"AAAAAAAAAAAAAAAAAAAAAAAAAA="),
      "keyId" : NumberLong(0)
    }
  }
}
```

Step 4. Create a 3-node replica set for config server.

4.1) On /etc/mongod.conf of each node, update the configuration.

- Change
storage:
dbPath: /data/db
- Add port: 27019
- Comment #bindIp line
- Add
replication:
replSetName: config_rs
- Add
sharding:
clusterRole: configsvr

where to write logging data.

```
systemLog:
  destination: file
  logAppend: true
  path: /var/log/mongodb/mongod.log
```

Where and how to store data.

```
storage:
  dbPath: /data/db
  journal:
    enabled: true
```

engine:

mmapv1:

wiredTiger:

how the process runs

```
processManagement:
  fork: true # fork and run in background
  pidFilePath: /var/run/mongodb/mongod.pid # location of pidfile
  timeZoneInfo: /usr/share/zoneinfo
```

network interfaces

```
net:
  port: 27019
  bindIp: 0.0.0.0 # Enter 0.0.0.0,:: to bind to all IPv4 and IPv6 addresses
```

#security:

#operationProfiling:

```
replication:
  replSetName: config_rs
```

```
sharding:
  clusterRole: configsvr
```

Enterprise-Only Options

#auditLog:

#snmp:

4.2) Start mongod using

```
sudo mongod --config /etc/mongod.conf
```

4.3) Connect a mongo shell to one of the mongod instances and run rs.initiate() per shard.

```
sudo mongo --port 27019
```

```
config = {_id : "config_rs",  
members:[{_id:0,host:"PRIVATE_IP_ADDRESS:27019"},  
          {_id:1,host:"PRIVATE_IP_ADDRESS: 27019"},  
          {_id:2,host:"PRIVATE_IP_ADDRESS: 27019"}]}}  
rs.initiate(config)
```

```
> config = {_id : "config_rs", members:[{_id:0,host:"172.31.31.46:27019"},{_id:1,host:"172.31.25.251:27019"},{_id:2,host:"172.31.30.219:27019"}]}}  
{  
  "_id" : "config_rs",  
  "members" : [  
    {  
      "_id" : 0,  
      "host" : "172.31.31.46:27019"  
    },  
    {  
      "_id" : 1,  
      "host" : "172.31.25.251:27019"  
    },  
    {  
      "_id" : 2,  
      "host" : "172.31.30.219:27019"  
    }  
  ]  
}  
|>  
|> rs.initiate(config)  
{  
  "ok" : 1,  
  "operationTime" : Timestamp(1546838072, 1),  
  "$gleStats" : {  
    "lastOpTime" : Timestamp(1546838072, 1),  
    "electionId" : ObjectId("00000000000000000000000000000000")  
  },  
  "lastCommittedOpTime" : Timestamp(0, 0),  
  "$clusterTime" : {  
    "clusterTime" : Timestamp(1546838072, 1),  
    "signature" : {  
      "hash" : BinData(0,"AAAAAAAAAAAAAAAAAAAAAAAAAAAA="),  
      "keyId" : NumberLong(0)  
    }  
  }  
}
```

Step 5. Configure mongos, add shards, and connect to clients.

5.1) On /etc/mongod.conf of each node, update the configuration.

- Comment the storage setting, as it is not going to store data.
- Comment #bindIp line.
- Add configDB under sharding.

sharding:

configDB: config_rs/Private_IP_Address:27019,
Private_IP_Address:27019,Private_IP_Address:27019

where to write logging data.

systemLog:

destination: file

logAppend: true

path: /var/log/mongodb/mongod.log

Where and how to store data.

#storage:

dbPath: /var/lib/mongo

journal:

enabled: true

engine:

mmapv1:

wiredTiger:

how the process runs

processManagement:

fork: true # fork and run in background

pidFilePath: /var/run/mongodb/mongod.pid # location of pidfile

timezoneInfo: /usr/share/zoneinfo

network interfaces

net:

port: 27017

bindIp: 0.0.0.0 # Enter 0.0.0.0,:: to bind to all IPv4 and IPv6 addresses or,

#security:

#operationProfiling:

#replication:

sharding:

configDB: config_rs/172.31.31.46:27019,172.31.25.251:27019,172.31.30.219:27019

Enterprise-Only Options

#auditLog:

5.2) Start mongos using

sudo mongos --config /etc/mongod.conf

5.3) Connect a mongo shell

sudo mongo --port 27017

5.4) Add two shards and their replica sets.

```
sh.addShard("rs1/PRIVATE_IP_ADDRESS:27018,PRIVATE_IP_ADDRESS:27018,PRIVATE_IP_ADDRESS:27018")
```

```
sh.addShard("rs2/PRIVATE_IP_ADDRESS:27018,PRIVATE_IP_ADDRESS:27018,PRIVATE_IP_ADDRESS:27018")
```

```
mongos> sh.addShard("rs1/172.31.30.123:27018,172.31.29.92:27018,172.31.30.237:27018")
```

```
{
  "shardAdded" : "rs1",
  "ok" : 1,
  "operationTime" : Timestamp(1546839444, 6),
  "$clusterTime" : {
    "clusterTime" : Timestamp(1546839444, 6),
    "signature" : {
      "hash" : BinData(0,"AAAAAAAAAAAAAAAAAAAAAAAAAAAA="),
      "keyId" : NumberLong(0)
    }
  }
}
```

```
mongos> sh.addShard("rs2/172.31.33.99:27018,172.31.45.208:27018,172.31.27.49:27018")
```

```
{
  "shardAdded" : "rs2",
  "ok" : 1,
  "operationTime" : Timestamp(1546839483, 5),
  "$clusterTime" : {
    "clusterTime" : Timestamp(1546839483, 6),
    "signature" : {
      "hash" : BinData(0,"AAAAAAAAAAAAAAAAAAAAAAAAAAAA="),
      "keyId" : NumberLong(0)
    }
  }
}
```

5.5) Enable sharding on the database and collection.

- Enable sharding for a database.

```
sh.enableSharding("mydb")
```

- Determine the shard key, create an index on the shard key and shard the collection.

```
use mydb
```

```
db.friend.createIndex( {"name":"hashed"} )
```

```
sh.shardCollection("mydb.friend", {"name":"hashed"})
```

Step 6. Load data from S3.

- configure ec2 using
aws configure
- import using aws s3 cp and mongoimport (should enable sharding on the db and collection first)
aws s3 cp s3://bucketname/file - | mongoimport --db
database_name(mydb) --collection collection_name(friend) --port
27017

```
[ec2-user@ip-172-31-17-100 ~]$ aws s3 cp s3://usfca-msan694/world_bank_project.json - | mongoimport --db msds697 --collection world_bank_project --port 27017
2019-01-07T07:30:37.912+0000    connected to: localhost:27017
2019-01-07T07:30:38.812+0000    num failures: 15
; "http://www.worldbank.org/projects/P131763/first-programmatic-development-policy-loan?lang=en" } does not contain shard key for pattern { borrower: "hashed"
" }
2019-01-07T07:30:38.814+0000    imported 485 documents
```

Reference :

1. Convert a Replica Set to a Sharded Cluster, MongoDB,
<https://docs.mongodb.com/manual/tutorial/convert-replica-set-to-replicated-shard-cluster/>
2. Install MongoDB Community Edition on Amazon Linux, MongoDB,
<https://docs.mongodb.com/manual/tutorial/install-mongodb-on-amazon/>
3. How to setup MongoDB Sharded Cluster with Replicasets on AWS, Cloud Buddy,
<https://m.youtube.com/watch?v=-sLpMqVBZI>